



Assignment 3 - Stencil



- ▶ The game of LIFE
 - ▶ 1970
 - ▶ A cellular automaton
 - 2D grid of cells
 - Two states: Dead, Alive
 - ▶ Each cell reproduces depending on its state and the state of its 8 neighbor cells
 - Born if exactly 3 neighbors alive
 - Survives if 2 or 3 neighbors alive
 - ▶ Turing complete!
 - Rendell, 2002
- ▶ Author: John Conway
 - ▶ Born Dec 26, 1937
 - ▶ Died Apr 11, 2020, COVID-19



▶ Rule 110

- ▶ 1985
- ▶ 1D grid of cells, two states
- ▶ One of 256 rules possible for 1D-neighborhood two-state cells

$a[i] = (110 \gg (4*a[i-1] + 2*a[i] + a[i+1])) \& 1;$


- ▶ Turing complete!
 - Cook, 2005

▶ Author: Stephen Wolfram

- ▶ Born Aug 29, 1959
- ▶ www.wolframalpha.com


- ▶ Many mathematical models work as follows:
 - ▶ A (multi-dimensional) array of cells
 - Each cell has a state
 - ▶ A global clock
 - In each tick, each cell is updated based on the (previous) states of its neighbors
 - The update formula is usually independent of position and time
- ▶ Examples:
 - ▶ Cellular automata
 - Ulam & von Neumann, 1948
 - Finite number of states of each cell
 - 1D: Rule 110; 2D: LIFE
 - ▶ Stencils
 - Emmons, 1944
 - Numeric solution of partial differential equations (finite difference method)
 - Heat transfer, hydrodynamics, ...
 - State of each cell consists of several real or complex numbers, depending on the problem

Serial evaluation of a stencil




- ▶ Two alternating buffers required
 - ▶ Stencils propagate state in both directions
 - ▶ The new states must always be computed from the old states
- ▶ Boundary cells may require special handling
 - ▶ Wrapping around is the simplest (but usually not physically correct)

Evaluation of a part of the space




- ▶ If the state is not known outside a part of the space, the result is valid only for a subset of the space
 - ▶ Shrinking with time

Parallel evaluation of a stencil




- ▶ A running thread cannot continuously observe the results produced by adjacent threads
 - ▶ Synchronization in every generation would be too expensive
- ▶ An individual task is to compute G generations, producing W results
 - ▶ The input to such a task is $G+W+G$ wide
 - ▶ A part of the work is duplicated in adjacent tasks (threads)

Parallel evaluation of a stencil




- ▶ An individual task is to compute G generations, producing W results
 - ▶ The input to such a task is $G+W+G$ wide
- ▶ The same thread may continue with subsequent generations
 - ▶ Synchronization with adjacent threads is required each G generations

Data exchanged between threads



- ▶ The same thread may continue with subsequent generations
 - ▶ Synchronization with adjacent threads is required each G generations
 - ▶ Two buffers of size G are sent from each thread to its neighbors
 - ▶ Two buffers of size G are received by each thread from its neighbors

Data exchanged between threads



- ▶ Every G generations, synchronization is required
 - ▶ Before synchronization, each thread holds W valid elements
 - Plus two buffers of size G , containing out-of-date elements
 - ▶ Two buffers of size G must be copied from each thread to its neighbors
- ▶ Two possible approaches
 - ▶ The receiving thread copies the data from adjacent threads to its own empty buffers
 - The adjacent threads may copy in parallel but they must not start computing further generations
 - Synchronization (rendezvous) required both before and after copying!
 - ▶ The sending thread copies the data and then sends the data including ownership of the buffer
 - The out-of-date buffers may be used, provided they are separable
 - Only one synchronization required (waiting for the messages)

► Synchronization primitives in C++17

- Starting and joining threads
 - Requires a master thread
 - Thread start is significantly slower than synchronization between running threads
- Mutex
 - Not suitable for waiting for an event
- Promise-future
 - Not reusable!
 - Additional thread-safe mechanism for dispatching promises (or futures) required
- Condition variable
 - Difficult to understand, coupled with a mutex
 - Spurious triggers
 - Reusable, does not involve allocation on each use

▶ Emulating rendezvous (between two threads)

- ▶ Corresponds to `std::barrier` in C++20
- ▶ Could be emulated by a pair of `std::binary_semaphore` in C++20
 - But you only have C++17
- ▶ With `std::condition_variable`:
 - Lock a mutex
 - Increment a shared counter
 - Notify a condition variable
 - While the counter is less than 2, wait for the condition variable
 - Waiting internally unlocks the mutex, allowing the other thread to increment
 - Spurious wakeups may occur – repeated testing of the counter is required
 - Unlock the mutex
 - BEWARE: A thread-safe way of resetting the counter is required
 - Resetting must appear after all the threads safely exited the rendezvous but before any of them enters the rendezvous again
 - In our case, we need a rendezvous both before and after the copying – resetting one of them may be done inside the other (if controlled by the same mutex)

► Emulating message passing

- Easily emulated by a `std::counting_semaphore` in C++20
 - But you only have C++17
- With `std::condition_variable`:
 - Send
 - With a mutex locked, push the message to a queue
 - Notify a condition variable
 - Receive
 - Lock the mutex
 - While the queue is empty, wait for the condition variable
 - Pop the message from the queue
 - Unlock the mutex
- In our case
 - Each pair of adjacent threads needs two message channels
 - Each message channel may contain up to two messages
 - A general (dynamically allocating) queue is not needed

The assignment

- ▶ Your job is to implement:
 - ▶ A generic structure used to hold the states of cells
 - 1-dimensional, wrapped-around = circle
 - ▶ A generic method which executes a given stencil function
 - For a given number of generations
 - In parallel, using only the C++17 standard library
 - ▶ Submit into recodex as "stencil1d.hpp"
- ▶ The testing framework will apply two stencil functions:
 - ▶ Rule 110
 - State of a cell = bool
 - ▶ Lemmings
 - State of a cell is a structure containing some small numbers
- ▶ The framework will test correctness by printing some states
 - ▶ Using recodex to compare against the desired output
- ▶ Beware: Recodex will also apply some not-so-infinite time limits
- ▶ The framework and the desired outputs are available outside Recodex
 - ▶ [//www.ksi.mff.cuni.cz/teaching/nprg051-web/DU/du-1920-3.framework.zip](http://www.ksi.mff.cuni.cz/teaching/nprg051-web/DU/du-1920-3.framework.zip)

The required interface

```
template< typename ET> class circle {
public:
    circle(std::size_t s);
    std::size_t size() const;
    void set(std::ptrdiff_t x, const ET& v);
    REF get(std::ptrdiff_t x) const;

    template< typename SF>
    void run(
        SF&& sf,
        std::size_t g,
        std::size_t thrs =
            std::thread::hardware_concurrency());
};
```

► Class `circle<ET>`

- **ET** is the type representing the status of one cell
- The constructor allocates the space for the desired number of cells
 - also returned by `size()`
 - each cell initialized as `ET()`
- The function **set/get** serve for writing/reading individual cells
 - The cells are circularly arranged – index overflows are handled by wrapping
 - The **set/get** function must support indexes (x) in the range `<-size(), 2*size()-1>`
 - BEWARE: the built-in operator% does not perform mathematical modulo for negative numbers
 - get may return by const reference but beware of `std::vector<bool>`

The required interface

```
template< typename ET> class circle {
public:

    circle(std::size_t s);

    std::size_t size() const;

    void set(std::ptrdiff_t x, const ET& v);

    REF get(std::ptrdiff_t x) const;

    template< typename SF>
    void run(
        SF&& sf,
        std::size_t g,
        std::size_t thrs =
            std::thread::hardware_concurrency());
};
```

- ▶ The function **run** does the parallel evaluation of the stencil
 - ▶ The stencil is specified as a function/functor/lambda with the following interface:
`ET sf(ET left, ET mine, ET right);`
 - The arguments may also be passed by const reference; therefore, the run function shall avoid copying the cell states when calling sf
 - ▶ The run function performs **g** generations, modifying the contents of the underlying class
 - Any auxiliary resources like threads or work buffers shall be freed upon return
 - ▶ The **thrs** argument specifies the desired number of threads, the default is as shown
 - Other parameters, like suitable values of W and G, must be determined internally
 - Beware of non-divisibility